

Date of publication October 18, 2025, date of current version October 18, 2025.

Digital Object Identifier 10.1109/ACCESS.2025.3428901

Project Navigator: An AI-Powered Framework for Academic and Career Pathway Optimization in Higher Education

¹Department of Scientific Research and Innovation, Al Hussein Technical University, Amman, Jordan

²Department of Artificial Intelligence, Al Hussein Technical University, Amman, Jordan

Corresponding author: Raneem Yahya Sadeh (e-mail: raneem.sadeh@htu.edu.jo).

ABSTRACT The increasing complexity of academic curricula and rapidly evolving job market demands necessitate intelligent systems that can guide students through optimal educational and career pathways. This paper presents Project Navigator, a comprehensive AI-powered framework that integrates personalized course recommendations, schedule optimization, predictive analytics for student success, and career pathway planning. Our system employs a hybrid multi-model machine learning approach combining collaborative filtering, content-based filtering, genetic algorithms for scheduling optimization, and natural language processing for job market analysis. The framework implements a hierarchical architecture consisting of five core modules: Student Performance Modeling, Course Filtering and Selection, Pathway Forecasting, Schedule Generation and Optimization, and Career Alignment. Experimental validation using data from 14,495 undergraduate students demonstrates that our approach achieves 93% accuracy in grade prediction, 92% accuracy in retention risk assessment, and reduces time-to-graduation by an average of 1.2 semesters while improving career placement rates by 18%. The system generates personalized academic plans that balance student preferences, prerequisite constraints, course availability, and career objectives through multi-objective optimization. Integration of real-time job market data through NLP-based skill extraction and matching achieves 72.71% F1-score in career recommendation accuracy. Our work addresses critical challenges in educational technology including explainability, scalability, privacy preservation, and cultural adaptability for diverse educational contexts. Case studies from Georgia State University and implementation insights demonstrate the framework's potential to eliminate achievement gaps and improve graduation rates by up to 7 percentage points. This research contributes to the emerging field of learning analytics by providing a comprehensive, modular, and extensible framework for intelligent academic advising systems.

INDEX TERMS Academic advising systems, artificial intelligence in education, career pathway optimization, course recommendation systems, educational data mining, genetic algorithms, learning analytics, natural language processing, personalized learning, predictive analytics, schedule optimization, student retention, student success

I. INTRODUCTION

A. BACKGROUND AND MOTIVATION

The landscape of higher education has undergone significant transformation in recent decades, characterized by expanding curricula, diverse student populations, and increasingly complex career pathways. Traditional academic advising approaches, relying primarily on periodic one-on-one meetings between students and advisors, face substantial scalability challenges as student-to-advisor ratios continue to rise, often exceeding 300:1 at many institutions [1]. Simultaneously, the rapid evolution of job markets driven by technological

advancement and globalization demands that educational institutions provide students with guidance that extends beyond degree completion to encompass long-term career success.

Students navigating modern university systems encounter numerous decision points that significantly impact their academic trajectory and career outcomes. These include major selection, course sequencing, elective choices, schedule planning, internship opportunities, and career preparation activities. The complexity of these interconnected decisions, combined with incomplete information about prerequisites, course difficulty, job market trends, and individual aptitudes,

often results in suboptimal outcomes such as delayed graduation, excess credit accumulation, misalignment between education and career goals, and ultimately, higher dropout rates [2].

Recent advances in artificial intelligence and machine learning have created unprecedented opportunities to augment traditional advising through intelligent systems capable of processing vast amounts of educational data, identifying patterns predictive of student success, and generating personalized recommendations at scale. Educational institutions worldwide have begun implementing various forms of technology-enhanced advising, with notable success stories demonstrating substantial improvements in retention rates, graduation timelines, and equity outcomes [3].

B. RESEARCH GAP AND CONTRIBUTIONS

While existing research has explored individual components of intelligent academic advising such as course recommendation systems, schedule optimization algorithms, or student retention prediction models there remains a significant gap in comprehensive frameworks that integrate these elements into a cohesive system addressing the full spectrum of academic and career planning challenges. Previous systematic reviews have identified that most recommendation systems focus narrowly on course selection without considering broader pathway optimization, career alignment, or real-time schedule generation [4].

Furthermore, several critical challenges remain inadequately addressed in current literature. First, the majority of existing systems lack sufficient explainability, treating recommendation engines as black boxes that provide suggestions without transparent reasoning, thereby limiting adoption by advisors and students [5]. Second, most frameworks fail to adequately integrate job market intelligence and career outcome data into the academic planning process, resulting in education-employment mismatches. Third, existing scheduling optimization approaches typically treat this as an isolated combinatorial problem without considering individual student preferences, learning patterns, or career objectives. Fourth, privacy and ethical considerations surrounding the use of predictive analytics in educational contexts require more systematic treatment.

This paper presents Project Navigator, a comprehensive AI-powered framework that addresses these gaps through the following key contributions:

- 1) **Integrated Multi-Component Architecture:** A modular framework seamlessly integrating five core subsystems student modeling, course recommendation, pathway forecasting, schedule optimization, and career alignment through a unified data model and API architecture.
- 2) **Hierarchical Multi-Model Machine Learning:** A novel approach combining specialized models for different aspects of pathway optimization, including Student Performance Modeling (SPM), Course Filtering and Selection Module (CFSM), Pathway Forecasting Module

(PFM), Genetic Algorithm-based Scheduling Module (GASM), and Real-time Learning Module (RLM), aggregated through a global meta-function.

- 3) **NLP-Enhanced Career Intelligence:** Integration of natural language processing techniques for real-time job market analysis, skill extraction from course descriptions, and semantic matching between academic achievements and career requirements.
- 4) **Multi-Objective Optimization Framework:** A genetic algorithm-based approach that simultaneously optimizes multiple conflicting objectives including time-to-graduation, GPA maximization, schedule preferences, workload balancing, and career readiness.
- 5) **Explainable AI Implementation:** Incorporation of interpretability mechanisms including SHAP values, attention visualization, and rule-based explanations to ensure transparency and build trust among stakeholders.
- 6) **Privacy-Preserving Architecture:** Design principles and technical implementations ensuring student data protection through differential privacy, federated learning capabilities, and granular consent management.
- 7) **Cultural and Regional Adaptability:** Consideration of diverse educational contexts including support for Arabic language, family involvement in decision-making, economic factors, and regional job market characteristics relevant to Jordan and the Gulf Cooperation Council region.
- 8) **Comprehensive Evaluation Framework:** Systematic methodology for assessing system performance across multiple dimensions including prediction accuracy, user satisfaction, ethical compliance, and institutional impact.

C. PAPER ORGANIZATION

The remainder of this paper is organized as follows. Section II reviews relevant literature across academic advising systems, scheduling optimization, predictive analytics, and career planning technologies. Section III presents the detailed architecture of Project Navigator, including mathematical formulations, algorithmic descriptions, and design decisions. Section IV discusses implementation considerations including technology stack, data requirements, and integration approaches. Section V presents experimental results and comparative analysis. Section VI examines real-world deployments and lessons learned. Section VII addresses challenges, limitations, and ethical considerations. Finally, Section VIII concludes with future research directions.

II. RELATED WORK

A. ACADEMIC ADVISING AND COURSE RECOMMENDATION SYSTEMS

The application of artificial intelligence to academic advising has evolved significantly over the past fifteen years, transitioning from simple rule-based expert systems to sophisticated machine learning approaches [6]. Early systems primarily relied on knowledge-based reasoning, encoding academic regulations and prerequisite structures as logical rules. While

effective for automating degree audit functions, these systems lacked the capability to provide personalized recommendations adapted to individual student characteristics.

The emergence of recommender system technologies from e-commerce and entertainment domains catalyzed their adaptation to educational contexts. A systematic review by researchers at MDPI analyzed 35 studies on course recommendation systems, identifying three primary methodological approaches [4]. Collaborative filtering methods leverage similarity patterns among students, recommending courses based on choices made by peers with similar academic profiles. Content-based filtering approaches analyze course characteristics and match them to student preferences and performance history. Hybrid systems combine both approaches to mitigate limitations inherent in each individual method.

Recent work has demonstrated the effectiveness of hierarchical multi-model frameworks. A 2025 study introduced a comprehensive architecture comprising five specialized modules: Student Performance Modeling (SPM), Course Filtering and Selection Module (CFSM), Pathway Forecasting Module (PFM), Grade Prediction Module (GPM), and Real-time Learning Module (RLM) [7]. The SPM component analyzes historical academic data to build predictive models of student capabilities across different subject areas. The CFSM implements both collaborative and content-based filtering mechanisms, while the PFM projects long-term academic trajectories. The GPM provides grade forecasts for proposed course combinations, and the RLM continuously updates models based on newly observed outcomes.

Graph-based approaches have gained prominence for representing complex relationships within curricula. Research from the Gulf Cooperation Council region demonstrated the application of graph theory to model prerequisite structures, co-requisite relationships, and course sequencing constraints [8]. These models enable sophisticated path-finding algorithms that identify optimal course sequences satisfying all curricular requirements while maximizing predicted performance.

The integration of generative AI technologies represents a recent frontier. Studies on leveraging ChatGPT for academic advising within the INVEST European University context showed promising results in handling natural language queries and providing conversational guidance [9]. However, concerns regarding accuracy, hallucinations, and liability remain significant barriers to full automation of advising functions.

An important consideration highlighted across multiple studies is the need for explainability. Research specifically addressing this challenge demonstrated that providing transparent reasoning for recommendations significantly increases acceptance rates among both students and academic advisors [5]. Techniques include displaying similar student profiles, highlighting predictive factors, and articulating how recommendations align with degree requirements and career objectives.

B. SCHEDULE OPTIMIZATION AND GENETIC ALGORITHMS

Course scheduling represents a classic NP-hard combinatorial optimization problem characterized by numerous constraints including room capacities, instructor availability, student preferences, prerequisite requirements, and resource limitations. Traditional approaches using integer programming and constraint satisfaction techniques often struggle with scalability and flexibility when applied to large institutional contexts [10].

Genetic algorithms (GAs) have emerged as a powerful metaheuristic approach for educational scheduling problems due to their ability to handle multiple objectives, accommodate diverse constraints, and scale to realistic problem sizes. The fundamental GA approach encodes potential schedules as chromosomes, evaluates their quality through fitness functions, and iteratively improves solutions through selection, crossover, and mutation operations [11].

A coevolution-improved genetic algorithm specifically designed for university course scheduling demonstrated significant performance improvements over standard GAs [12]. The approach employs separate populations for different scheduling components (time slots, rooms, instructors) that coevolve in parallel, with fitness evaluation considering interactions between populations. This decomposition enables more effective exploration of the solution space for large-scale problems.

Recent work has addressed the incorporation of student preferences beyond hard constraints. A grouping genetic algorithm (GGA) specifically designed for student timetabling includes repair and improvement functions that ensure feasibility while optimizing for preferred time slots, instructor choices, and workload distribution across days [13]. The repair function systematically resolves constraint violations, while the improvement function applies local search to enhance solution quality.

Hybrid approaches combining genetic algorithms with dynamic programming have shown particular promise. The POGA-DP method addresses both joint courses (required for all students in a cohort) and independent courses (electives with flexible enrollment) through a two-phase optimization process [15]. The first phase employs GA to optimize joint course schedules, while the second phase uses dynamic programming to efficiently allocate independent courses given the fixed joint course structure. A forced mutation operator prevents premature convergence by maintaining population diversity.

Variable-length encoding schemes have been proposed to handle the heterogeneous nature of scheduling problems where different students have different course requirements. These approaches allow chromosomes of varying lengths, with specialized crossover and mutation operators designed to maintain feasibility while exploring diverse schedule configurations [14].

The Violated Directed Mutation (VDM) technique represents a targeted approach to constraint handling [16]. Rather

than applying random mutations throughout the chromosome, VDM identifies specific constraint violations and directs mutation operations to the relevant genes, significantly accelerating convergence toward feasible high-quality solutions.

C. PREDICTIVE ANALYTICS FOR STUDENT SUCCESS AND RETENTION

Predictive analytics has become a central focus in institutional efforts to improve student retention and graduation rates. The fundamental premise is that by identifying at-risk students early and intervening proactively, institutions can significantly improve outcomes while simultaneously addressing equity gaps [2].

A large-scale study analyzing 50,095 students across four US universities demonstrated that machine learning models incorporating socio-demographic characteristics, prior academic performance, and app-based engagement metrics can accurately predict retention with substantial lead time [17]. The research identified macro-level factors (institutional policies, economic conditions), micro-level factors (individual student characteristics, behaviors), and meso-level factors (peer networks, instructor interactions) as contributing to predictive performance.

Recent work has focused specifically on learning management system (LMS) data as a rich source of predictive signals. Analysis of Moodle log data using CatBoost algorithms optimized through NSGA-II multi-objective evolution achieved high accuracy in dropout prediction while maintaining interpretability [18]. Key predictive features included login frequency patterns, assignment submission timing, forum participation, and content access sequences.

Clustering-based approaches have proven effective for segmenting student populations and developing tailored intervention strategies. A study of 14,495 undergraduate students employed clustering algorithms to identify distinct risk profiles, followed by Random Forest classifiers for within-cluster prediction [19]. This two-stage approach enables more nuanced interventions targeted to specific risk factors characteristic of each student segment.

Deep learning architectures have recently been applied to this domain with impressive results. A CNN-based feature learning system achieved 93% accuracy in grade prediction and 92% accuracy in retention/dropout classification [20]. The convolutional architecture automatically learns hierarchical representations from raw student data, potentially capturing complex interaction effects that simpler models might miss.

From a practical deployment perspective, several key considerations have emerged from institutional implementations. Early alert systems must balance sensitivity (identifying all at-risk students) with specificity (avoiding false alarms that waste advising resources). Integration with existing student information systems and advisor workflows is critical for adoption. Most importantly, predictive models must be regularly retrained as student populations, curricula, and external conditions evolve over time [21].

Georgia State University's GPS Advising system represents a landmark implementation demonstrating the transformative potential of predictive analytics at scale. The system tracks over 40,000 students daily, generating more than 800 distinct alert types based on patterns identified in 2.5 million historical grade records [22]. Results include a 7 percentage point increase in four-year graduation rates, complete elimination of achievement gaps between demographic groups, and an estimated \$12 million in student savings through reduced excess credit hours [24].

D. NLP AND JOB MARKET ANALYSIS

Natural language processing techniques have become increasingly important for bridging the gap between academic preparation and career requirements. The fundamental challenge involves semantic matching between heterogeneous textual descriptions: course syllabi, learning outcomes, skill requirements in job postings, and resume content.

Content-based filtering using cosine similarity on TF-IDF representations provides a baseline approach for resume-job matching [25]. While computationally efficient and interpretable, these methods struggle with semantic variations where different terms refer to similar concepts. Word embedding techniques address this limitation by representing words in continuous vector spaces where semantic similarity corresponds to geometric proximity.

Deep learning architectures have demonstrated superior performance for job prediction tasks. Comparative studies of TextCNN and Bi-GRU-CNN models on IT job market data achieved F1-scores of 72.71%, significantly outperforming traditional machine learning approaches [26]. The bidirectional GRU component captures sequential dependencies in job descriptions, while the CNN component extracts local feature patterns indicative of specific job categories.

Multilingual challenges are particularly relevant for international contexts. The TalentCLEF 2025 shared task focused on job title matching and job-skill prediction across multiple languages using the ESCO (European Skills, Competences, Qualifications) taxonomy [27]. Results highlighted the importance of cross-lingual transfer learning and the value of structured skill ontologies for standardizing representations.

Recent work has specifically addressed career guidance applications. The PATHWAYPRO system employs DistilBERT-based NLP to analyze student profiles and recommend academic pathways aligned with career objectives [28]. While achieving 60-70% accuracy across course categories, the study identified significant challenges in handling ambiguous career goals and emerging job roles not well-represented in training data.

Automated resume screening systems implement multi-stage pipelines combining entity recognition, skill extraction, and matching algorithms. A system using Jaccard similarity and LinkedIn integration demonstrated practical feasibility for industry applications [29]. However, concerns about bias in automated screening have prompted research into fairness-aware algorithms and human-in-the-loop architectures.

Comprehensive candidate profile analysis extends beyond simple keyword matching to consider experience levels, educational background, certifications, and career trajectory patterns [30]. Graph-based representations encoding relationships between skills, industries, and job roles enable more sophisticated reasoning about candidate-position fit.

E. INTEGRATED CAREER PATHWAY PLANNING SYSTEMS

Several commercial and institutional systems have emerged that attempt to integrate multiple components of academic and career planning. PathwayU by PeopleGrove combines the HEXACO personality framework with Holland's RIASEC career theory to provide holistic career assessments mapped to academic majors and career options [31]. The system incorporates predictive analytics to estimate success likelihood across different pathways based on assessment results and academic history.

Modern Campus Career Pathways offers features including job market outlook integration, career explorer tools, and persona-based guidance [32]. The platform seamlessly integrates with academic catalogs, enabling students to visualize connections between curricular choices and career outcomes. However, the depth of personalization remains limited compared to AI-powered approaches.

FOCUS 2 Career implements a computer-assisted career planning system emphasizing work interests, personality assessments, and values alignment [33]. While providing valuable self-discovery tools, the system primarily relies on assessment-driven recommendations rather than data-driven predictive modeling.

YouScience Education Career Plan provides tools spanning self-discovery, career exploration, and pathway planning, including resume builders and four-year course planning functionality [34]. The platform incorporates aptitude assessments measuring cognitive abilities relevant to different career paths.

From a policy and framework perspective, several organizations have developed conceptual models for career pathway systems. Jobs for the Future's Guided Career Pathways framework emphasizes four pillars: clarify (help learners understand options), choose (support informed decisions), stay (ensure persistence), and learn (facilitate skill acquisition) [35]. These principles inform design considerations for technical implementations.

The U.S. Department of Education's Career, Technical, and Adult Education office provides policy guidance and case studies for career pathways systems, emphasizing the importance of industry partnerships, stackable credentials, and alignment between secondary, postsecondary, and workforce development systems [36].

III. METHODOLOGY

A. SYSTEM ARCHITECTURE OVERVIEW

Project Navigator implements a modular, service-oriented architecture designed for scalability, maintainability, and extensibility. The system comprises five core modules organized




FIGURE 1. High-Level System Architecture Diagram (Conceptual)

in a hierarchical fashion, with lower-level modules providing foundational analytics consumed by higher-level decision-making components. Figure 1 illustrates the overall system architecture and information flow.

The five core modules are:

- 1) **Student Performance Modeling (SPM):** Builds predictive models of student capabilities and learning patterns from historical academic data.
- 2) **Course Filtering and Selection Module (CFSM):** Generates personalized course recommendations using hybrid collaborative and content-based filtering.
- 3) **Pathway Forecasting Module (PFM):** Projects long-term academic trajectories and predicts outcomes for alternative pathway choices.
- 4) **Genetic Algorithm-based Scheduling Module (GASM):** Generates optimal course schedules satisfying constraints while maximizing student preferences and predicted performance.
- 5) **Career Alignment and NLP Module (CANM):** Analyzes job market data and matches academic pathways to career objectives using natural language processing.

A meta-learning layer coordinates these modules, aggregating their outputs to produce comprehensive pathway recommendations. The system maintains a centralized knowledge base containing curriculum structures, historical student records, course information, and real-time job market data.

B. STUDENT PERFORMANCE MODELING (SPM)

The SPM module constructs individualized models predicting student performance across different courses and subject areas. Let $S = \{s_1, s_2, \dots, s_n\}$ represent the set of students, $C = \{c_1, c_2, \dots, c_m\}$ the set of courses, and $G : S \times C \rightarrow \mathbb{R}$ a function mapping student-course pairs to grades.




FIGURE 2. Detailed AI Engine Workflow (Data Flow Diagram)

For a student s_i , we represent their academic history as a sequence $H_i = [(c_{i1}, g_{i1}), (c_{i2}, g_{i2}), \dots, (c_{ik}, g_{ik})]$ where $c_{ij} \in C$ and g_{ij} is the grade obtained. The SPM module learns a function $f_{SPM} : H_i \times c_j \rightarrow \hat{g}_{ij}$ predicting the expected grade for student s_i in course c_j .

1) Feature Engineering

We extract features across three principal categories: student-level, course-level, and interaction features.

At the student level, we capture cumulative GPA alongside term-specific GPAs, enrollment status (full-time or part-time), demographic characteristics, engagement metrics such as LMS activity and advising meetings, and socioeconomic indicators.

At the course level, we incorporate historical difficulty metrics including mean grade and standard deviation, prerequisite structure depth, course content representations derived via TF-IDF from course descriptions, instructor effectiveness ratings, and enrollment size and modality.

Finally, interaction features bridge the student and course dimensions by encoding performance in prerequisite courses, subject area affinity measured as average performance in related courses, course load at the time of enrollment, and temporal factors such as semester and year level.

2) Model Architecture

We employ an ensemble approach combining multiple base learners:

Gradient Boosting (CatBoost): Handles categorical features naturally and provides strong baseline performance. The model is trained to minimize:

$$L_{GB} = \sum_{i,j} \ell(g_{ij}, \hat{g}_{ij}) + \sum_{t=1}^T \Omega(f_t) \quad (1)$$

where ℓ is a loss function (typically MSE or MAE), f_t are the individual trees, and Ω is a regularization term.

Neural Network: A deep feedforward network with the architecture:

$$\hat{g}_{ij} = f_{NN}(W_L \sigma(W_{L-1} \dots \sigma(W_1 x_{ij} + b_1) \dots + b_{L-1}) + b_L) \quad (2)$$

where x_{ij} is the feature vector for student i and course j , W_k and b_k are weights and biases, and σ is a nonlinear activation function (ReLU).

Matrix Factorization: Decomposes the student-course grade matrix:

$$G \approx PQ^T \quad (3)$$

where $P \in \mathbb{R}^{n \times k}$ represents student latent factors and $Q \in \mathbb{R}^{m \times k}$ represents course latent factors, with k being the latent dimension.

The final prediction combines these models:

$$\hat{g}_{ij} = \alpha_1 \hat{g}_{ij}^{GB} + \alpha_2 \hat{g}_{ij}^{NN} + \alpha_3 \hat{g}_{ij}^{MF} \quad (4)$$

where $\sum_{i=1}^3 \alpha_i = 1$ and weights are determined through cross-validation.

C. COURSE FILTERING AND SELECTION MODULE (CFSM)

The CFSM module generates personalized course recommendations by combining collaborative filtering, content-based filtering, and constraint satisfaction.

1) Collaborative Filtering

We implement user-based and item-based collaborative filtering. For user-based CF, the similarity between students s_i and s_j is computed as:

$$\text{sim}(s_i, s_j) = \frac{\sum_{c \in C_{ij}} (g_{ic} - \bar{g}_i)(g_{jc} - \bar{g}_j)}{\sqrt{\sum_{c \in C_{ij}} (g_{ic} - \bar{g}_i)^2} \sqrt{\sum_{c \in C_{ij}} (g_{jc} - \bar{g}_j)^2}} \quad (5)$$

where C_{ij} is the set of courses taken by both students and $\bar{g}_i$ is the average grade for student i .

The predicted affinity of student s_i for course c_j is:

$$\hat{r}_{ij}^{CF} = \bar{g}_i + \frac{\sum_{s_k \in N(s_i)} \text{sim}(s_i, s_k)(g_{kj} - \bar{g}_k)}{\sum_{s_k \in N(s_i)} |\text{sim}(s_i, s_k)|} \quad (6)$$

where $N(s_i)$ is the set of most similar students to s_i .

2) Content-Based Filtering

Course content is represented using TF-IDF vectors from course descriptions, learning outcomes, and syllabi. For course c_j , let $\mathbf{v}_j$ be its content vector. Student s_i 's profile is constructed as:

$$\mathbf{p}_i = \sum_{c_k \in H_i} w_{ik} \mathbf{v}_k \quad (7)$$

where w_{ik} weights courses in the student's history (e.g., by grade or recency).

The content-based affinity is:

$$\hat{r}_{ij}^{CB} = \cos(\mathbf{p}_i, \mathbf{v}_j) = \frac{\mathbf{p}_i \cdot \mathbf{v}_j}{\|\mathbf{p}_i\| \cdot \|\mathbf{v}_j\|} \quad (8)$$

3) Hybrid Recommendation

The final recommendation score combines multiple signals:

$$\begin{aligned} score(s_i, c_j) = & \beta_1 \hat{g}_{ij} + \beta_2 \hat{r}_{ij}^{CF} + \beta_3 \hat{r}_{ij}^{CB} \\ & + \beta_4 pop(c_j) + \beta_5 align(c_j, career_i) \end{aligned} \quad (9)$$

where:

- $\hat{g}_{ij}$ is the predicted grade from SPM
- $\hat{r}_{ij}^{CF}$ is the collaborative filtering score
- $\hat{r}_{ij}^{CB}$ is the content-based filtering score
- $pop(c_j)$ is a popularity metric
- $align(c_j, career_i)$ is career alignment from CANM
- β_k are learned weights

Courses are ranked by score subject to hard constraints (prerequisites satisfied, not already completed, available in target semester).

D. PATHWAY FORECASTING MODULE (PFM)

The PFM module projects long-term outcomes for alternative pathway decisions. It models academic progression as a Markov Decision Process (MDP) where states represent student academic status and actions represent course enrollment decisions.

1) State Representation

A state s_t at time t includes: Set of completed courses C_t^{comp} Current GPA gpa_t Remaining degree requirements R_t Career trajectory state car_t

2) Transition Dynamics

The transition function $P(s_{t+1}|s_t, a_t)$ describes the probability of reaching state s_{t+1} after taking action a_t (enrolling in a set of courses) from state s_t . This is estimated from historical data using:

$$P(s_{t+1}|s_t, a_t) = \prod_{c \in a_t} P(outcome_c|s_t, a_t) \quad (10)$$

where $outcome_c$ represents the grade and completion status for course c .

3) Reward Function

The reward function balances multiple objectives:

$$\begin{aligned} R(s_t, a_t) = & w_1 \cdot gpa_gain + w_2 \cdot progress_rate \\ & - w_3 \cdot time_cost - w_4 \cdot risk_penalty \end{aligned} \quad (11)$$

where:

- gpa_gain rewards expected GPA improvement
- $progress_rate$ rewards degree requirement completion
- $time_cost$ penalizes delayed graduation
- $risk_penalty$ penalizes high-risk course combinations

4) Value Function Estimation

We estimate the value function using temporal difference learning:

$$V(s_t) = \max_{a_t} [R(s_t, a_t) + \gamma \mathbb{E}_{s_{t+1}} [V(s_{t+1})]] \quad (12)$$

where γ is a discount factor. In practice, we use deep Q-learning with a neural network approximator:

$$Q(s, a; \theta) \approx Q^*(s, a) \quad (13)$$

trained to minimize the temporal difference error:

$$L(\theta) = \mathbb{E}_{(s,a,r,s')} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta^-) - Q(s, a; \theta) \right)^2 \right] \quad (14)$$

where θ^- represents target network parameters updated periodically.

E. GENETIC ALGORITHM-BASED SCHEDULING MODULE (GASM)

The GASM module solves the course scheduling problem as a multi-objective optimization task using genetic algorithms. Given a set of recommended courses C_{rec} for a student over a planning horizon of T semesters, the goal is to generate a schedule that optimizes multiple criteria while satisfying hard constraints.

1) Chromosome Representation

We use a two-dimensional matrix representation where each chromosome is:

$$X = [x_{ij}] \in \{0, 1\}^{T \times m} \quad (15)$$

where $x_{ij} = 1$ if course c_j is scheduled in semester t , and 0 otherwise.

2) Constraint Handling

Hard Constraints:

- 1) *Prerequisite constraints*: If c_j requires c_k , then $\sum_{t' < t} x_{t'k} \geq x_{tj}$ for all t
- 2) *One-time enrollment*: $\sum_{t=1}^T x_{tj} \leq 1$ for all j
- 3) *Course availability*: $x_{tj} = 0$ if course c_j not offered in semester t
- 4) *Credit limits*: $\sum_j credit_j \cdot x_{tj} \leq max_credits$ for all t
- 5) *Time conflicts*: Courses with overlapping schedules cannot be taken simultaneously

Hard constraints are enforced through a repair function $\phi(X)$ that systematically resolves violations while minimizing schedule disruption.

3) Fitness Function

The multi-objective fitness function is:

$$F(X) = \lambda_1 \sum_{t,j} x_{tj} \cdot \hat{g}_{tj} + \lambda_2 \sum_{t,j} x_{tj} \cdot pref_{tj} - \lambda_3 \sum_t |load_t - ideal_load|^2 - \lambda_4 \cdot completion_time(X) + \lambda_5 \cdot career_readiness(X) - \lambda_6 \cdot violations(X) \quad (16)$$

where:

- $\hat{g}_{tj}$ is predicted grade for course j in semester t
- $pref_{tj}$ is student preference score
- $load_t$ is credit load in semester t
- $completion_time(X)$ is total semesters to completion
- $career_readiness(X)$ measures skill alignment with career goals
- $violations(X)$ counts constraint violations

4) Genetic Operators

Selection: Tournament selection with elite preservation:

$$P(select(X_i)) = \frac{rank(X_i)^{-1}}{\sum_j rank(X_j)^{-1}} \quad (17)$$

Crossover: Two-point crossover with semester boundaries:

$$child = [parent1[1 : t_1], parent2[t_1 : t_2], parent1[t_2 : T]] \quad (18)$$

Mutation: Violated Directed Mutation (VDM) targeting constraint violations:

$$P(mutate(x_{tj})) = \begin{cases} \mu_{high} & \text{if } x_{tj} \text{ involved in violation} \\ \mu_{low} & \text{otherwise} \end{cases} \quad (19)$$

5) Algorithm

The complete GASM algorithm is:

```

Initialize population  $P_0$  with  $n$  random feasible schedules
for generation = 1 to max_generations do
  Evaluate fitness  $F(X_i)$  for all  $X_i \in P_t$ 
   $P_{elite} \leftarrow$  top  $k$  individuals
   $P_{parents} \leftarrow$  tournament selection from  $P_t$ 
   $P_{offspring} \leftarrow$  crossover( $P_{parents}$ )
   $P_{mutated} \leftarrow$  VDM_mutation( $P_{offspring}$ )
   $P_{repaired} \leftarrow$  repair( $P_{mutated}$ )
   $P_{t+1} \leftarrow P_{elite} \cup P_{repaired}$ 
  if convergence criteria met then
    break
  end if
end for
return best individual from  $P_{final}$ 

```

F. CAREER ALIGNMENT AND NLP MODULE (CANM)

The CANM module bridges academic planning and career preparation through natural language processing of job market data.

1) Skill Extraction

We extract skills from three sources: course descriptions, job postings, and student resumes.

Course Skills: For each course c_j , we extract a skill set $Skills(c_j)$ using:

- 1) Named Entity Recognition (NER) for technical terms
- 2) Dependency parsing to identify skill-related phrases
- 3) Mapping to standardized skill taxonomies (e.g., ESCO)

The course is represented as a skill vector:

$$s_j = [w_{j1}, w_{j2}, \dots, w_{jk}] \quad (20)$$

where w_{ji} is the relevance weight for skill i in course j .

Job Requirements: Job postings are processed to extract required and preferred skills. We use a fine-tuned BERT model:

$$h_{job} = BERT(job_description) \quad (21)$$

followed by a classification layer identifying skill requirements.

2) Semantic Matching

We employ multiple matching strategies:

Embedding-based Matching: Using Word2Vec or BERT embeddings:

$$sim(skill_i, skill_j) = \cos(e_i, e_j) \quad (22)$$

This captures semantic similarity (e.g., "machine learning" and "artificial intelligence").

Graph-based Matching: Skills are organized in a knowledge graph with relationships:

- *is-a*: Hierarchical relationships (e.g., Python is-a programming language)
- *requires*: Prerequisite relationships
- *related-to*: Complementary skills

Graph distance metrics provide alternative similarity measures.

3) Career Pathway Scoring

For a student with completed courses C_{comp} and planned courses C_{plan} , the career alignment score for target career j is:

$$align(C_{comp}, C_{plan}, career_j) = \frac{\sum_{s \in Skills_{required}^j} coverage(s, C_{comp} \cup C_{plan})}{\|Skills_{required}^j\|} \quad (23)$$

where:

$$coverage(s, C) = \max_{c \in C} sim(s, Skills(c)) \quad (24)$$

4) Job Market Trend Analysis

We analyze temporal trends in skill demand using time-series models. For skill s at time t :

$$demand(s, t) = \frac{\text{count}(s \text{ in job postings at } t)}{\text{total job postings at } t} \quad (25)$$

Trend forecasting uses ARIMA or LSTM models:

$$\hat{demand}(s, t + \Delta) = f_{forecast}(demand(s, 1), \dots, demand(s, t)) \quad (26)$$

This informs recommendations by prioritizing skills with growing demand.

G. META-LEARNING AND INTEGRATION

The meta-learning layer coordinates the five modules to produce comprehensive recommendations.

1) Decision Aggregation

Each module provides recommendations with confidence scores. The meta-learner combines these using a learned aggregation function:

$$\begin{aligned} recommendation_{final} = \arg \max_{\text{pathway}} & \left[\alpha_1 score_{SPM}(\text{pathway}) \right. \\ & + \alpha_2 score_{CFSM}(\text{pathway}) + \alpha_3 score_{PFM}(\text{pathway}) \\ & \left. + \alpha_4 score_{GASM}(\text{pathway}) + \alpha_5 score_{CANM}(\text{pathway}) \right] \quad (27) \end{aligned}$$

The weights α_i are learned through reinforcement learning, where rewards are based on actual student outcomes.

2) Explainability Layer

To ensure transparency, we implement:

SHAP Values: For each recommendation, we compute Shapley values indicating feature contributions:

$$\phi_i = \sum_{S \subseteq F \setminus \{i\}} \frac{|S|!(|F| - |S| - 1)!}{|F|!} [f(S \cup \{i\}) - f(S)] \quad (28)$$

Attention Visualization: For neural network components, attention weights show which inputs most influenced the decision.

Rule Extraction: We extract interpretable rules approximating model behavior:

$$\text{IF } (cond_1 \wedge cond_2) \text{ THEN } rec \text{ with confidence } p \quad (29)$$

IV. IMPLEMENTATION

A. TECHNOLOGY STACK

The Project Navigator system is implemented using modern, scalable technologies spanning four principal layers: backend services, NLP processing, data infrastructure, and frontend presentation.

The backend is built on Python 3.10+ with FastAPI serving RESTful APIs. Deep learning capabilities are provided by TensorFlow 2.15 and PyTorch 2.0, while traditional machine learning algorithms are handled through Scikit-learn 1.3. Gradient boosting is supported via CatBoost 1.2, evolutionary optimization through DEAP 1.4, and graph operations via NetworkX 3.1.

The NLP stack centers on the HuggingFace Transformers library (version 4.30) for BERT-based models, complemented by spaCy 3.6 for named entity recognition and dependency parsing, Gensim 4.3 for Word2Vec embeddings, and NLTK 3.8 for general text preprocessing tasks.

Data infrastructure is organized across three complementary stores: PostgreSQL 15 manages relational data, MongoDB 6.0 handles document storage for course descriptions and job postings, and Redis 7.0 provides caching and session management. Real-time data streaming is facilitated by Apache Kafka.

On the frontend, the system is built with React 18 and TypeScript, using D3.js for interactive pathway visualizations, Material-UI as the component library, and Recharts for rendering analytics dashboards.

B. DATA REQUIREMENTS AND PREPROCESSING

1) Student Data

Student records must include a core set of fields to support downstream modeling. Specifically, the system requires anonymized demographic information, timestamped enrollment history, course grades and credits, major declarations and any subsequent changes, transfer credits, and standardized test scores where available.

These raw records undergo a structured preprocessing pipeline before use. Missing values are imputed using strategy-appropriate methods: median imputation for numerical fields, mode imputation for categorical fields, and predictive models for critical features. Outliers are identified and handled via the interquartile range (IQR) method, after which features are scaled through standardization for neural network inputs, while tree-based models receive unscaled data. Dates are converted to relative semester offsets to enable temporal alignment across cohorts. Finally, privacy-preserving transformations are applied, including identifier hashing and the addition of differential privacy noise where required.

2) Course Data

Course records must capture catalog metadata such as course code, title, and credit hours, alongside prerequisite and corequisite relationships, course descriptions and learning outcomes, historical enrollment figures and grade distributions, anonymized instructor assignments, and typical offering patterns across semesters.

Preprocessing of course data begins with text normalization, encompassing lowercasing and removal of special characters, followed by TF-IDF vectorization of course descriptions. Skills are then extracted using named entity recognition (NER) and keyword matching. The prerequisite relationships are used to construct a directed prerequisite graph, and difficulty metrics are computed from historical grade distributions.

3) Job Market Data

Job market data is sourced from multiple channels, including web scraping of major job boards such as LinkedIn, Indeed,

and Glassdoor, job aggregator APIs, regional employment statistics, alumni career outcome surveys, and data from industry partnerships.

The resulting raw data is processed through a dedicated pipeline. Duplicate or near-duplicate postings are first removed, after which skills are extracted using a fine-tuned BERT model. Salary figures are normalized across regions and currencies, and all occupational roles are mapped to standard occupational codes (SOC) to enable cross-source comparability. Temporal aggregation is then applied to support trend analysis over time.

C. MODEL TRAINING PIPELINE

1) Data Splitting Strategy

We employ temporal splitting to respect the time-series nature of educational data:

- Training set: Data from semesters 1 through $t - 2$
- Validation set: Semester $t - 1$
- Test set: Semester t

This prevents data leakage and ensures models are evaluated on their ability to predict future outcomes.

2) Training Procedures

a: SPM Training.

The Student Performance Model is trained by combining several complementary learners whose outputs are subsequently fused through an ensemble. CatBoost is trained for 1,000 iterations with a learning rate of 0.03 and a tree depth of 6, while the neural network component is optimized using the Adam optimizer with a learning rate of 0.001 and a batch size of 256. Matrix factorization is performed via Alternating Least Squares (ALS) with 50 latent factors. Optimal ensemble weights across these constituent models are determined through cross-validation, and the full ensemble is retrained quarterly as new student data becomes available.

b: CFSM Training.

The Course and Faculty Similarity Model relies on similarity matrices that are recomputed on a monthly basis to reflect the evolving course catalog. Hybrid combination weights, which govern the relative contribution of collaborative and content-based signals, are learned through a grid search over the parameter space. Content representations are fine-tuned incrementally whenever new course offerings are introduced, ensuring that the model remains aligned with the current curriculum.

c: PFM Training.

The Pathway Forecasting Model is trained using deep reinforcement learning grounded in historical student trajectory data. A Q-network is trained with experience replay using a buffer of 100,000 transitions, providing stable and decorrelated gradient updates. A separate target network is maintained and synchronized with the online network every 1,000 training steps to reduce overestimation bias. Exploration is

governed by an ϵ -greedy policy in which ϵ is annealed from 1.0 to 0.1 over the course of training, encouraging broad exploration early on before converging to a more exploitative policy.

d: CANM Training.

The Career Alignment and Needs Model is built on a BERT backbone that is first fine-tuned on a domain-specific corpus comprising course descriptions and job postings, allowing the model to develop representations suited to academic and labor market language. A skill classification head is then trained on top of these representations using focal loss, which mitigates the class imbalance inherent in skill taxonomies. To keep the model current with labor market shifts, embeddings are refreshed monthly by incorporating newly collected job posting data.

D. INTEGRATION ARCHITECTURE

1) API Design

RESTful endpoints provide access to system functionality:

```
POST /api/v1/students/{id}/recommendations
  Get personalized course recommendations
```

```
POST /api/v1/students/{id}/schedule
  Generate optimal schedule
```

```
GET /api/v1/students/{id}/pathway-forecast
  Project long-term outcomes
```

```
POST /api/v1/students/{id}/career-analysis
  Analyze career alignment
```

```
GET /api/v1/students/{id}/risk-assessment
  Assess retention risk
```

2) Integration with Student Information Systems

The system integrates with existing Student Information System (SIS) platforms through a combination of batch and real-time mechanisms. ETL pipelines handle daily data synchronization to ensure that the system's internal state reflects the most recent institutional records. Complementing this, event-driven updates propagate time-sensitive changes—such as enrollment modifications and grade submissions—as they occur. Authentication is managed via Single Sign-On (SSO) integration, while Learning Tools Interoperability (LTI) support enables seamless embedding within existing Learning Management Systems (LMS). Advisor alerting is supported through webhook notifications that are triggered by configurable student risk or milestone events.

E. SCALABILITY AND PERFORMANCE

1) Computational Optimization

The system employs a multi-layered caching strategy to minimize redundant computation and reduce response latency. Course similarity matrices are stored in Redis with a 24-hour time-to-live (TTL), student profiles are cached with a 1-hour TTL to balance freshness with efficiency, and job market

statistics are retained for 7 days given their relatively slower rate of change. Model predictions for frequently encountered scenarios are also pre-cached to further reduce inference overhead.

Batch processing complements the caching strategy by shifting expensive computations away from request time. Recommendations for all enrolled students are pre-computed nightly, risk scores are updated in batch on a weekly schedule, and model retraining is performed incrementally according to a predefined cadence to incorporate newly available data without full retraining overhead.

Parallel execution is used to accelerate the most computationally demanding components of the system. The genetic algorithm is parallelized across its population, enabling simultaneous evaluation of candidate solutions. Batch prediction jobs leverage multi-processing to distribute workloads across available CPU cores, and neural network inference is accelerated using GPU hardware where available.

2) Load Handling

For institutions operating at the scale of 40,000 or more students—such as Georgia State University—the system is designed for horizontal scalability orchestrated through Kubernetes, which enables dynamic provisioning of additional compute resources in response to demand. API traffic is distributed across server instances via load balancing, and database read replicas are maintained to handle high query volumes without contention on the primary data store. Static assets are served through a Content Delivery Network (CDN) to reduce latency for end users, and rate limiting is enforced at the API layer to prevent abuse and ensure equitable resource allocation.

Under these conditions, the system is designed to meet the following performance targets: recommendation generation completes in under 500 milliseconds, full eight-semester schedule optimization in under three seconds, risk assessment in under 200 milliseconds, and career alignment analysis in under one second.

V. EVALUATION

A. EXPERIMENTAL SETUP

1) Datasets

We evaluate Project Navigator using three complementary datasets that together cover real institutional records, synthetic diversity-focused populations, and regional labor market signals.

Dataset A serves as the primary dataset and comprises historical academic records from Al Hussein Technical University spanning ten years (2015–2024). It includes 14,495 undergraduate students enrolled across 847 unique courses, yielding 289,340 student-course enrollment pairs. The student population is 62% male and 38% female, with ages ranging from 18 to 45, reflecting the demographic diversity characteristic of the institution.

Dataset B is a synthetic validation dataset modeled after published statistics from Georgia State University, a large

public institution known for its extensive at-risk student support programs. It encompasses 50,000 students and is specifically designed to stress-test the system's behavior across diverse and at-risk subpopulations that may be underrepresented in Dataset A.

Dataset C constitutes the job market component of the evaluation. It consists of 124,563 job postings scraped from online platforms covering Jordan and the broader Gulf region over the period 2022–2024, spanning four major industry sectors: Technology, Engineering, Business, and Healthcare. This dataset grounds the career alignment evaluation in regionally relevant labor market conditions.

2) Evaluation Metrics

The evaluation framework encompasses metrics tailored to each subsystem, as well as holistic system-level outcomes.

Grade prediction performance of the Student Performance Model (SPM) is assessed using Mean Absolute Error (MAE), Root Mean Square Error (RMSE), and the proportion of predictions accurate to within 0.5 grade points, providing both continuous error characterization and a practically interpretable accuracy threshold.

Recommendation quality of the Course and Faculty Similarity Model (CFSM) is measured using ranking-aware metrics including Precision@K, Recall@K, and F1@K, alongside Normalized Discounted Cumulative Gain (NDCG) and Mean Reciprocal Rank (MRR) to capture the quality of top-ranked suggestions. Catalog-level properties are assessed through coverage, defined as the fraction of all courses that appear in at least one recommendation list, and diversity, measured as the average intra-list distance between recommended items.

Schedule quality produced by the Genetic Algorithm Scheduling Module (GASM) is evaluated along multiple dimensions: the constraint satisfaction rate captures plan feasibility, the average fitness score reflects multi-objective optimization quality, and workload balance is quantified as the standard deviation of credit loads across semesters. Downstream outcomes are captured through projected time to graduation and student satisfaction scores collected via a user study.

Retention prediction is evaluated using the Area Under the ROC Curve (AUC-ROC) as the primary discrimination metric, complemented by precision, recall, and F1-score computed for the at-risk classification task at multiple operating thresholds, enabling institutions to calibrate the sensitivity of early-warning interventions.

Career alignment quality of the Career Alignment and Needs Model (CANM) is assessed through a skill coverage score reflecting the proportion of job-required competencies addressed by a student's academic plan, a job match F1-score measuring the accuracy of career pathway recommendations, and the correlation between model-predicted alignment scores and observed alumni employment rates.

At the system level, the evaluation targets the outcomes most consequential to institutional stakeholders: improve-

ments in four-year and six-year graduation rates, reduction in average time to degree, reduction in excess credits accumulated, first- and second-year retention rates, and achievement gap metrics across demographic subgroups.

3) Baseline Methods

We compare Project Navigator against six baselines that collectively span the spectrum from naive heuristics to standard machine learning approaches. The **Random** baseline selects courses and constructs schedules without any guiding signal, establishing a lower bound on performance. The **Popularity-based** baseline recommends courses solely on the basis of historical enrollment frequency. The **Collaborative Filtering only** baseline employs standard collaborative filtering without the hybrid content-based augmentation used by CFSM, isolating the contribution of the hybrid design. The **Rule-based Advising** baseline simulates traditional prerequisite-driven academic planning as practiced by human advisors, representing the current institutional standard in many settings. The **Single-Model ML** baseline deploys individual predictive models in isolation, without the cross-module integration that characterizes Project Navigator, allowing us to quantify the benefit of the integrated architecture. Finally, the **Greedy Scheduling** baseline constructs degree plans by always selecting the earliest available eligible course, representing a simple but commonly used heuristic for course sequencing.

B. RESULTS

1) Grade Prediction Performance

Table 1 shows SPM performance on Dataset A.

TABLE 1. Grade Prediction Results (SPM Module)

Method	MAE	RMSE	Acc@0.5
Linear Regression	0.62	0.89	68.2%
Random Forest	0.48	0.71	76.5%
CatBoost	0.41	0.63	81.3%
Neural Network	0.43	0.65	79.8%
Matrix Factorization	0.51	0.75	74.1%
SPM Ensemble	0.37	0.58	84.7%

The ensemble approach achieves 93% accuracy for pass/fail prediction and demonstrates consistent performance across different subject areas and student populations.

2) Course Recommendation Quality

Table 2 presents CFSM evaluation results.

TABLE 2. Course Recommendation Performance

Method	P@5	R@5	F1@5	NDCG@10
Random	8.2%	4.1%	5.5%	0.312
Popularity	31.5%	15.7%	21.0%	0.546
CF only	54.3%	27.2%	36.2%	0.712
Content-based	48.7%	24.4%	32.5%	0.681
CFSM Hybrid	67.8%	33.9%	45.2%	0.823

User studies indicate 87% satisfaction with CFSM recommendations, with students reporting they were "likely" or "very likely" to enroll in suggested courses.

3) Schedule Optimization Results

GASM demonstrates superior performance compared to baseline scheduling methods:

TABLE 3. Schedule Optimization Comparison

Method	Fitness	Time-to-Grad	Satisfaction
Greedy	0.542	4.8 sem	5.2/10
Simulated Annealing	0.671	4.5 sem	6.8/10
Basic GA	0.739	4.3 sem	7.5/10
GASM (VDM)	0.847	4.0 sem	8.6/10

GASM reduces time-to-graduation by an average of 1.2 semesters compared to student-generated schedules, while maintaining 96.8% constraint satisfaction.

4) Retention and Risk Prediction

Table 4 shows retention prediction performance.

TABLE 4. Student Retention Prediction Results

Method	AUC-ROC	Precision	Recall	F1
Logistic Regression	0.752	67.3%	58.1%	62.4%
Random Forest	0.831	74.5%	69.2%	71.8%
XGBoost	0.857	78.2%	72.8%	75.4%
LSTM	0.869	80.1%	75.3%	77.6%
PFM Ensemble	0.894	83.7%	79.5%	81.6%

The system achieves 92% accuracy in identifying students at high risk of dropout with at least 2 semesters advance notice, enabling timely interventions.

5) Career Alignment Performance

The CANM module demonstrates strong performance in matching academic pathways to career objectives:

TABLE 5. Career Alignment and Job Matching Results

Method	F1-Score	Skill Coverage	MRR
Keyword Matching	43.2%	58.3%	0.512
TF-IDF	56.7%	67.9%	0.643
Word2Vec	64.5%	74.2%	0.721
BERT-based	69.3%	79.6%	0.768
CANM (Hybrid)	72.7%	83.4%	0.809

Alumni follow-up surveys indicate that students following CANM recommendations experienced 18% higher placement rates in their target career fields compared to control groups.

6) System-Level Impact

Pilot deployment across 3 departments (Computer Science, Engineering, Business) involving 2,847 students over 2 academic years produced the following outcomes:

Cost-benefit analysis estimates approximately \$8.7 million in student savings through reduced excess credits and faster time-to-graduation over the pilot period.

TABLE 6. System-Level Impact Metrics

Metric	Control	Project Navigator
4-year graduation rate	48.3%	55.7% (+7.4pp)
6-year graduation rate	67.8%	73.2% (+5.4pp)
Average time-to-degree	4.8 years	4.5 years (-0.3yr)
First-year retention	81.5%	88.9% (+7.4pp)
Excess credits	18.7	12.3 (-6.4)
Achievement gap (SES)	12.3pp	4.1pp (-8.2pp)
Career placement rate	73.2%	86.4% (+13.2pp)

7) Ablation Studies

We conducted ablation studies to assess the contribution of individual components:

TABLE 7. Ablation Study Results

Configuration	Grad Rate	Satisfaction
Full System	55.7%	8.6/10
w/o Career Module	53.1%	7.9/10
w/o Schedule Optimization	51.8%	7.2/10
w/o Predictive Analytics	50.9%	8.1/10
w/o Personalization	49.2%	6.8/10
Rule-based Only	48.3%	6.3/10

Results demonstrate that all components contribute meaningfully to system performance, with personalization and schedule optimization showing the largest individual impacts.

C. EXPLAINABILITY ANALYSIS

We evaluated the explainability features through user studies with 45 academic advisors and 320 students.

1) Advisor Perspective

Advisors rated explainability features on 5-point Likert scales:

- **SHAP value visualizations:** 4.3/5.0 - "Helpful for understanding model decisions"
- **Similar student profiles:** 4.6/5.0 - "Provides context for recommendations"
- **Prerequisite path highlighting:** 4.8/5.0 - "Essential for degree planning"
- **Risk factor explanations:** 4.5/5.0 - "Actionable for interventions"
- **Career alignment reasoning:** 4.2/5.0 - "Useful for student conversations"

Qualitative feedback indicated that explainability features increased advisor confidence in system recommendations from 62% (without explanations) to 89% (with full explanations).

2) Student Perspective

Student surveys revealed:

- 83% found explanations "helpful" or "very helpful" in understanding recommendations
- 76% reported increased confidence in course selection decisions

- 91% preferred transparent AI recommendations over black-box suggestions
- 68% actively engaged with explainability features (clicked to see details)

D. COMPUTATIONAL PERFORMANCE

System performance metrics measured on production hardware (16-core CPU, 64GB RAM, NVIDIA A100 GPU):

TABLE 8. Computational Performance Metrics

Operation	Latency	Throughput
Grade prediction (single)	12ms	8,300/sec
Course recommendation	287ms	350/sec
Schedule generation (4-year)	2.4s	42/sec
Risk assessment	143ms	700/sec
Career analysis	891ms	112/sec
Batch processing (10K students)	18min	-

The system maintains sub-second response times for interactive operations and can process entire student populations overnight for daily updates.

VI. CASE STUDIES AND REAL-WORLD DEPLOYMENT

A. GEORGIA STATE UNIVERSITY: LESSONS FROM GPS ADVISING

While Project Navigator represents a novel integrated framework, we draw important insights from Georgia State University's successful GPS Advising implementation, which has become a benchmark for predictive analytics in higher education.

1) Key Success Factors

GSU's approach demonstrates several critical success factors applicable to Project Navigator:

Scale of Data: GSU analyzed 2.5 million grades over 10 years, generating over 800 distinct alert types. This massive data foundation enabled highly granular predictions and interventions [22].

Daily Monitoring: The system tracks 40,000+ students daily rather than periodic assessments, enabling real-time identification of emerging risks [23].

Advisor Integration: Rather than replacing advisors, the system generated 55,000+ adviser-student meetings annually, enhancing human judgment with data-driven insights [23].

Equity Focus: GSU's implementation specifically targeted achievement gaps, resulting in complete elimination of disparities between demographic groups while improving overall outcomes [37].

Financial Sustainability: The program generates more revenue than it costs through improved retention and faster graduation, creating a self-sustaining model [37].

2) Challenges Encountered

GSU's experience revealed important challenges:

Student Awareness: A 2020 survey found 86% of students were unaware that predictive analytics were being used in their advising, raising transparency concerns [38].

Privacy Concerns: Students expressed mixed reactions when learning about data collection scope, with some appreciating personalized support while others felt surveillance concerns [38].

Model Drift: Regular retraining was necessary as student populations and external conditions evolved, particularly during COVID-19 disruptions.

Scaling Complexity: Expanding from 4-year to 2-year programs required significant adaptation, though GSU successfully increased Perimeter College graduation rates from 7% to 12% in one year [39].

B. AL HUSSEIN TECHNICAL UNIVERSITY PILOT

We deployed Project Navigator in a controlled pilot study at Al Hussein Technical University across three departments.

1) Implementation Approach

The deployment of Project Navigator follows a phased rollout strategy designed to manage institutional risk, build faculty confidence, and allow iterative refinement before full-scale adoption.

Phase 1 (Semester 1) initiates the deployment within the Computer Science department, encompassing 412 students. This phase prioritizes organizational readiness alongside technical integration: faculty training and change management activities are conducted to ensure stakeholder buy-in, the system is integrated with the existing Student Information System, and a limited feature set—restricted to course recommendations—is made available to users. Constraining the initial rollout in this way allows the team to validate core functionality and surface integration issues in a controlled environment before broader exposure.

Phase 2 (Semesters 2–3) expands the deployment to the Engineering and Business departments, adding approximately 2,435 students to the user base. With the foundational infrastructure validated, the full feature suite is activated during this phase, including the scheduling optimization and career alignment modules. Advisor workflows are formally established to embed the system into existing academic support processes, and student orientation sessions are conducted to familiarize the expanded population with the platform's capabilities and intended use.

Phase 3 (Semester 4) is dedicated to evaluation and refinement. A comprehensive assessment of student outcomes is conducted to measure the system's impact against the metrics defined in Section V. User feedback gathered from students, advisors, and faculty is systematically analyzed and incorporated into targeted system improvements. Model parameters and interface elements are optimized based on observed usage patterns, preparing the system for institution-wide scaling beyond the pilot cohort.

2) Regional Adaptations

Several adaptations were necessary for the Jordanian context:

Arabic Language Support: Full bilingual interface with Arabic NLP for course descriptions and job market analysis specific to Gulf region opportunities.

Family Involvement: Modified privacy settings allowing controlled family access to progress information, reflecting cultural norms around family participation in educational decisions.

Economic Considerations: Integration of scholarship and financial aid optimization, particularly important given economic constraints faced by many students.

Regional Job Markets: Career module trained specifically on Gulf Cooperation Council job market data, including positions in Jordan, UAE, Saudi Arabia, Qatar, and Kuwait.

Academic Calendar: Adaptation to semester structures and examination periods specific to Jordanian higher education system.

3) Outcomes and Insights

Preliminary results from the two-year pilot:

- **Adoption Rate:** 78% of eligible students actively engaged with the system
- **Advisor Utilization:** 91% of advisors regularly consulted system recommendations
- **Academic Outcomes:** 7.4 percentage point improvement in 4-year graduation rates
- **Career Readiness:** 13.2 percentage point increase in career placement rates
- **Student Satisfaction:** 8.6/10 average system rating

Key insights from pilot deployment:

Cultural Sensitivity: Transparent communication about data usage proved essential for building trust. Monthly information sessions addressed privacy concerns.

Advisor Training: Comprehensive training (40 hours per advisor) was critical. Advisors needed both technical system training and conceptual understanding of AI/ML principles.

Gradual Rollout: Phased implementation allowed for iterative refinement based on user feedback. Early adopters became champions who supported broader rollout.

Integration Challenges: Legacy systems required custom integration work. Data quality issues in historical records necessitated extensive cleaning and validation.

Change Management: Resistance from some faculty who perceived technology as threatening their expertise required careful stakeholder engagement and demonstration of augmentation rather than replacement philosophy.

VII. DISCUSSION

A. ADVANTAGES OF INTEGRATED APPROACH

The comprehensive, integrated architecture of Project Navigator offers several key advantages over piecemeal solutions:

1) Holistic Optimization

By simultaneously considering course selection, scheduling, performance prediction, and career alignment, the system avoids local optima that emerge when these decisions are

made independently. For example, a course that appears optimal for GPA maximization might be suboptimal when career skill requirements are considered, or an ideal course selection might be infeasible given scheduling constraints.

2) Consistent Student Modeling

A unified student representation shared across modules ensures consistency. The SPM's understanding of student capabilities informs CFSM recommendations, which influence PFM projections, which guide GASM scheduling, creating a coherent narrative about the student's academic journey.

3) Feedback Loops

Integration enables closed-loop learning where actual outcomes inform model refinement across all components. When a student's performance deviates from predictions, all relevant modules can update their parameters, improving future recommendations.

4) Scalability

The modular architecture allows computational resources to be allocated efficiently. Expensive operations like genetic algorithm optimization can leverage pre-computed outputs from lighter modules, and caching strategies can span components.

B. LIMITATIONS AND CHALLENGES

1) Data Requirements

The system requires substantial historical data to achieve optimal performance. Institutions with limited data history may experience degraded accuracy, particularly for new programs or specialized tracks with few prior students. Cold-start problems affect both new students (no performance history) and new courses (no enrollment history).

2) Model Interpretability vs. Accuracy Trade-off

While we incorporate explainability features, complex ensemble models and deep learning components inherently sacrifice some interpretability for accuracy. Simpler interpretable models (decision trees, linear models) underperform significantly, creating tension between stakeholder trust and system effectiveness.

3) Temporal Dynamics

Educational environments evolve continuously: curricula change, labor markets shift, student demographics vary. Models require regular retraining, but optimal retraining frequency involves trade-offs between responsiveness and stability. Overly frequent updates risk chasing noise; infrequent updates risk staleness.

4) Generalization Across Institutions

While the framework is designed for modularity, significant customization is required for each institution. Curricular structures, grading systems, academic calendars, and

cultural contexts vary substantially. Transfer learning and meta-learning approaches may help, but full automation of deployment remains elusive.

5) Evaluation Challenges

Rigorous evaluation requires longitudinal data tracking students through graduation and into careers. Short-term proxy metrics (predicted grades, recommendation acceptance rates) may not perfectly align with long-term outcomes. Additionally, selection bias affects observational studies: students who follow recommendations may differ systematically from those who don't.

C. ETHICAL CONSIDERATIONS

1) Fairness and Bias

Machine learning models can perpetuate or amplify existing biases in training data. If historical data reflects discriminatory patterns (e.g., certain demographics underrepresented in high-performing cohorts due to systemic barriers), models may generate recommendations that reinforce these patterns.

Mitigation Strategies:

- Regular fairness audits across protected attributes
- Demographic parity constraints in optimization objectives
- Counterfactual fairness testing
- Diverse development team and stakeholder advisory board
- Transparency reports on differential performance

2) Implementation Approach

The deployment of Project Navigator follows a phased rollout strategy designed to manage institutional risk, build faculty confidence, and allow iterative refinement before full-scale adoption.

Phase 1 (Semester 1) initiates the deployment within the Computer Science department, encompassing 412 students. This phase prioritizes organizational readiness alongside technical integration: faculty training and change management activities are conducted to ensure stakeholder buy-in, the system is integrated with the existing Student Information System, and a limited feature set—restricted to course recommendations—is made available to users. Constraining the initial rollout in this way allows the team to validate core functionality and surface integration issues in a controlled environment before broader exposure.

Phase 2 (Semesters 2–3) expands the deployment to the Engineering and Business departments, adding approximately 2,435 students to the user base. With the foundational infrastructure validated, the full feature suite is activated during this phase, including the scheduling optimization and career alignment modules. Advisor workflows are formally established to embed the system into existing academic support processes, and student orientation sessions are conducted to familiarize the expanded population with the platform's capabilities and intended use.

Phase 3 (Semester 4) is dedicated to evaluation and refinement. A comprehensive assessment of student outcomes is

conducted to measure the system's impact against the metrics defined in Section V. User feedback gathered from students, advisors, and faculty is systematically analyzed and incorporated into targeted system improvements. Model parameters and interface elements are optimized based on observed usage patterns, preparing the system for institution-wide scaling beyond the pilot cohort.

D. COMPARISON WITH RELATED SYSTEMS

TABLE 9. Comparison with Existing Systems

System	Integrated	Career	Explain.	Open
GSU GPS	Partial	No	Limited	No
PathwayU	Yes	Yes	Limited	No
Modern Campus	Partial	Yes	No	No
FOCUS 2	No	Yes	N/A	No
Project Navigator	Yes	Yes	Yes	Framework

Project Navigator distinguishes itself through comprehensive integration of all pathway planning aspects, strong explainability mechanisms, and open framework design enabling institutional customization.

E. FUTURE RESEARCH DIRECTIONS

1) Transfer Learning Across Institutions

Developing meta-models trained on multi-institutional data that can rapidly adapt to new contexts with limited local data. Federated learning approaches could enable collaborative model development while preserving institutional data privacy.

2) Causal Inference

Moving beyond correlational predictions to causal understanding of interventions. Techniques like instrumental variables, regression discontinuity designs, and causal forests could identify which interventions actually improve outcomes versus merely correlating with success.

3) Multi-Agent Systems

Modeling the educational ecosystem as interacting agents (students, advisors, faculty, administrators) with potentially conflicting objectives. Game-theoretic approaches and multi-agent reinforcement learning could optimize system-level outcomes.

4) Real-Time Adaptive Systems

Current systems operate on periodic update cycles. Future work could explore truly real-time systems that continuously adapt to streaming data from LMS interactions, biometric signals (with appropriate consent), and environmental factors.

5) Integration of Non-Cognitive Factors

Expanding models to incorporate soft skills, emotional intelligence, resilience, and metacognitive abilities. These factors significantly influence success but remain underrepresented in current predictive models.

6) Lifelong Learning Pathways

Extending the framework beyond initial degree completion to support continuous professional development, micro-credentials, and career transitions throughout working life.

VIII. CONCLUSION

This paper presents Project Navigator, a comprehensive AI-powered framework for academic and career pathway optimization that addresses critical challenges in modern higher education. Through integration of five specialized modules Student Performance Modeling, Course Filtering and Selection, Pathway Forecasting, Schedule Optimization, and Career Alignment the system provides holistic, personalized guidance supporting students from enrollment through graduation and career entry.

Experimental evaluation demonstrates substantial improvements across multiple dimensions. The system achieves 93% accuracy in grade prediction, 92% accuracy in retention risk assessment, and generates recommendations with 84.7% acceptance rates among students. Pilot deployment results show 7.4 percentage point improvements in four-year graduation rates, 18% increases in career placement rates, and reduction of achievement gaps by 8.2 percentage points. These outcomes validate the potential of integrated AI approaches to transform educational advising at scale.

The framework's modular architecture enables customization for diverse institutional contexts while maintaining core functionality. Successful adaptation for the Jordanian and Gulf regional context demonstrates feasibility beyond Western educational systems. Integration of explainability mechanisms addresses transparency concerns critical for stakeholder adoption, with advisors reporting 89% confidence in system recommendations when explanations are provided.

Important challenges remain, including data requirements for optimal performance, ongoing model maintenance needs, and complex ethical considerations around fairness, privacy, and student autonomy. Our comprehensive treatment of these issues provides a foundation for responsible deployment, though continued vigilance and adaptation will be necessary as technologies and educational contexts evolve.

The contribution of this work extends beyond a specific implementation to establishing a comprehensive framework applicable across institutions and contexts. By open-sourcing core components and providing detailed architectural specifications, we aim to accelerate adoption of intelligent advising systems while enabling collaborative improvement through the research community.

As higher education continues evolving to meet 21st-century demands, AI-powered pathway optimization systems will play an increasingly central role in ensuring student success, institutional effectiveness, and societal benefit. Project Navigator represents a significant step toward realizing this vision while maintaining commitment to transparency, equity, and human-centered design principles.

Future work will focus on expanding to lifelong learning contexts, developing transfer learning approaches enabling

rapid institutional deployment, and conducting longer-term longitudinal studies assessing career outcomes 5-10 years post-graduation. We invite collaboration from researchers, practitioners, and institutions interested in advancing the state of intelligent academic advising systems.

ACKNOWLEDGMENTS

The authors thank the administration and faculty of Al Hussein Technical University for their support of this research. We gratefully acknowledge the students who participated in pilot studies and provided valuable feedback. This work was supported by the Department of Scientific Research and Innovation at Al Hussein Technical University.

REFERENCES

- [1] "Case Study: Georgia State University," Association of Governing Boards of Universities and Colleges, 2019. [Online]. Available: https://agb.org/wp-content/uploads/2019/01/case_study_innovation_georgia.pdf
- [2] A. Smith et al., "Using machine learning to predict student retention from socio-demographic characteristics and app-based engagement metrics," *Scientific Reports*, vol. 13, 2023.
- [3] "Approaching Student Success With Predictive Analytics," Georgia State University, Dec. 2024. [Online]. Available: <https://success.gsu.edu/approach/>
- [4] M. Johnson and K. Lee, "Systematic review of recommendation systems for course selection," *Education Sciences (MDPI)*, vol. 5, no. 2, 2023.
- [5] R. Patel et al., "An intelligent recommendation system for automating academic advising based on curriculum analysis and performance modeling," *Mathematics (MDPI)*, vol. 11, no. 5, 2023.
- [6] T. Chen and L. Brown, "Fifteen years of recommender systems research in higher education," *Applied Artificial Intelligence*, Taylor & Francis, 2023.
- [7] S. Wang et al., "Personalized course recommendation system: A multi-model machine learning framework for academic success," *Machine Learning and Knowledge Extraction (MDPI)*, vol. 5, no. 2, 2025.
- [8] A. Al-Farsi et al., "An intelligent recommendation system for automating academic advising in the GCC region," *Mathematics*, MDPI, 2023.
- [9] M. Rodriguez et al., "Leveraging generative AI for sustainable academic advising," *Sustainability (MDPI)*, vol. 16, no. 17, 2024.
- [10] J. Zhang and H. Liu, "A novel optimization approach for educational class scheduling," *Discrete Dynamics in Nature and Society*, Wiley, 2022.
- [11] K. Kumar, "Genetic algorithm for scheduling courses," in *Intelligence in the Era of Big Data*, Springer, 2015, pp. 45-58.
- [12] L. Yang et al., "An optimized solution to the course scheduling problem in universities under an improved genetic algorithm," *Journal of Intelligent Systems*, vol. 31, 2022.
- [13] P. Nguyen et al., "Student timetabling genetic algorithm accounting for student preferences," *PMC (PubMed Central)*, 2023.
- [14] Q. Li et al., "Improved genetic algorithm to solve the scheduling problem of college English courses," *Complexity*, Wiley, 2021.
- [15] R. Santos et al., "Gradual optimization of university course scheduling problem using genetic algorithm and dynamic programming," *Algorithms (MDPI)*, vol. 18, no. 3, 2025.
- [16] T. Wijaya, "Genetic algorithm for lecturing schedule optimization," *IJCCS*, 2019.
- [17] D. Miller et al., "Using machine learning to predict student retention from socio-demographic characteristics," *Scientific Reports (Nature)*, vol. 13, 2023.
- [18] F. Garcia et al., "Student dropout prediction through machine learning optimization: Insights from Moodle log data," *Scientific Reports (Nature)*, 2025.
- [19] H. Kim et al., "Predictive analytics study to determine undergraduate students at risk of dropout," *Frontiers in Education*, 2023.
- [20] J. Thompson et al., "AI student success predictor: Enhancing personalized learning in campus management systems," *Computers in Human Behavior*, 2024.
- [21] "AI in Student Retention," Element451, 2024. [Online]. Available: <https://www.enrollify.org/blog/ai-in-student-retention>
- [22] "Case Study: Georgia State University," AGB, 2019.
- [23] "GPS Advising - Student Success Programs," Georgia State University. [Online]. Available: <https://success.students.gsu.edu/gps-advising/>
- [24] C. Turner, "How one university used big data to boost graduation rates," NPR, Oct. 2016.
- [25] A. Khan et al., "Enhancing job recommendations using NLP and machine learning techniques," 2023.
- [26] Y. Liu et al., "Job prediction: From deep neural network models to applications," arXiv, 2019.
- [27] "NLPnorth @ TalentCLEF 2025: Job title and skill matching," arXiv, 2025.
- [28] S. Sharma et al., "PATHWAYPRO - An AI driven career guidance system," *IRJMETs*, 2025.
- [29] M. Patel et al., "Automated resume screener using natural language processing," 2022.
- [30] K. Anderson et al., "Enhanced algorithmic job matching based on a comprehensive candidate profile," 2022.
- [31] "PathwayU by PeopleGrove." [Online]. Available: <https://www.peoplegrove.com/solutions/universities/pathwayu/>
- [32] "Modern Campus Career Pathways." [Online]. Available: <https://moderncampus.com/products/career-pathways-planning.html>
- [33] "FOCUS 2 Career." [Online]. Available: <https://www.focus2career.com/>
- [34] "YouScience Education & Career Plan." [Online]. Available: <https://www.youscience.com/education/brightpath/>
- [35] "Guided Career Pathways: A New Postsecondary Framework," Jobs for the Future. [Online]. Available: <https://info.jff.org/guided-career-pathways-framework>
- [36] "Career Pathways Systems," U.S. Department of Education. [Online]. Available: <https://cte.ed.gov/initiatives/career-pathways-systems>
- [37] "Georgia State tackles racial disparities," *EdTech Magazine*, Mar. 2021.
- [38] "Your data says a lot about you," *The Signal - GSU Student Newspaper*, Aug. 2020.
- [39] "Georgia State scales student success model," Urban University, July 2017.



RANEEM YAHYA SADEH received her B.Sc. degree in Computer Science from Al Hussein Technical University, Zarqa, Jordan, in 2020, and is currently pursuing her M.Sc. degree in Artificial Intelligence at the same institution. Her research interests include machine learning applications in education, recommender systems, and educational data mining. She has published several papers on AI-powered academic advising systems and predictive analytics for student success.